

Contents

Front Matter	2
Preface - JDD Briefly	2
Start	4
Execute	4
Automate	4
Review	5
Ensure	5
Demonstrate	5
Scale	5
Reinvent	6
Dedication	6
Errata	6
Acknowledgements	6
Process	7
Journey-Driven Agility	7
Engineering	8
Design	9
Initiation	9
DevOps	11
Architecture	13
API Forward	13
Forward From Behind the CMS	13
Forward of Systems of Record	14
Forward in Time	14
Uses Customer Language	15
API Forward and Microservices	15
Content Rocket	15
Data	16
Strategy	16
Experience Strategy	16
Content Strategy	16
Asset Creation	16
Activation	16
Journey Audience Map	17
Evolution	17
Reuse	17
Refine	17
Reevaluate	17
Techniques	17
Journey-Driven API Design	17
Document the decisions	17

Architecture Decision Record	18
Resource Analysis	18
1. Gather Functional Requirements	18
2. Gather Non Functional Requirements	18
3. Create a Domain Model (Conceptual ER Model)	18
4. Design an API Contract	18
System Analysis	18
Persistence Analysis	19
Maturing the Practice	20
Directive vs Suggestive Approaches	20
Example: Increase Unit Testing Adoption	21
Evolve Our Practice	21
Evolve Our Client's Practice to Include Unit Testing	22
Pattern Languages	24
People	24
Journey of the Software Professional	24
Manager Tools Adoption	24
Evolution of Teamwork	24
People Matter	24
Tools	25
Future Topics	25
Agile Release Planning	25
Code Review	26
Pull Request	26
Pairing as Review	26
Pride and Shame	26
Like, Wish, Wonder	26
Turn Watermelons Into Orangeberries	26
Back Matter	27
Bibliography	27

Front Matter

Preface - JDD Briefly

Journey-Driven Development (JDD) is how we create software systems for clients at VML. The phrase, 'for clients,' is important to us as it reminds us that our clients judge our work not only by how it runs but also by examining its construction critically as it is developed for hire.

We wrote JDD for ourselves, the developers, analysts, testers, architects, and operators of the software we create for our clients. We share it as we believe the

lessons learned are useful to software teams generally. You don't have to code to benefit from JDD but the subject matter is technical.

We work at an agency, VML. Sometimes we say advertising agency and sometimes we say marketing agency. We think of ourselves as agents of change and transformation.

Our software is delivered quickly. Our clients are demanding. We often guide the software through the initial launch and sometimes through years of updates, and our clients own the software we create and operate it long-term.

Goals:

1. Make high quality software fast.
2. Make it easy to learn our process.
3. Make it easy to modify our process.

We discovered JDD and we teach it to our clients as it is easy to adopt in parts or as a holistic transformation. We've applied JDD for large-scale transformation efforts in automotive, telecommunications, travel, finance, transportation, and construction.

JDD begins with API Forward, which simplifies complex systems integration. API Forward is event-driven, focusing on the events arising from the user journey. We map these events through lightweight orchestration via a forward facing facade. We implement the facade in a messaging system. That keeps the complexity of transactions away from our experience platform. See API Forward topic for details.

For the experience platform, JDD uses a technique we call Content Rocket. Content Rocket also begins with the user journey, specifically the audiences and moments discovered by our Content Systems practice. We map these to reusable content components that we can combine in different ways to both allow for specialization and maximize reuse. We've mapped Content Rocket to all of the major experience platforms. See Content Rocket topic for details.

None of this is completely foreign. We are informed and influenced by the conversations in our software development community. These conversations happen at conferences, in chat rooms, on blogs, and in books like this one. Here are some of our influences:

1. Agile [@rasmusson_agile_2014] generally and eXtreme Programming [@beck_extreme_2000] in particular
2. Software Development as Craft [@cunningham_software_2005]
3. DevOps [@kim_devops_2016]
4. Automation
5. Responsibility Driven Design (RDD)
6. Domain Driven Design (DDD)
7. Behavior Driven Development (BDD)
8. Test Driven Development (TDD)

Perhaps the first written version of Journey-Driven Development began as a one-page brief when a team that had been high-performing began to stumble. The memo served as a starting point for many other teams to define their process.

What follows is one fork of that document:

Start

The team begins when a technical lead and a project manager estimate the effort and staff the team.

- PM holds a dev kickoff. Each team member reads the scope, plan, and aligns their task with the larger ask.
- Everyone knows who the tech lead is.
- PM and tech lead are careful when they have to move resources mid-project and account for ramp time and communication overhead.
- PM and tech lead ensure team and individual accountability for hitting dates and budget. Be creative and escalate as needed.
- The tech lead validates dev team progress, setting estimates (and resetting remaining work as needed).

Execute

The team uses a project page or wiki to ensure up-to-date versions of these key artifacts are available to the entire team.

- Scope
- Project plan (or backlog), release plan, test plan, and defect list
- Requirements
- Wireframes, creative, and an Executable Style Guide.

Leader portion of weekly One-on-one Meeting references one or more of these documents (particularly the project plan).

The entire team:

- Attends stand-up (mandatory for dev team and PM).
- Refers to the above in their daily stand up using measurement language (e.g., “Yesterday I reviewed 3 requirements and I plan to finish 80% of 1 today”).
- Ensures the plan is realistic based on the estimated backlog and team velocity .

Automate

- Automate builds
- Automate tests - you don’t need to test everything, but complex areas should have regression tests.
- Automate deployment

Review

Review regularly. Pairing is faster than working alone and improves initial quality. Regular pairing gives us a structure for regular and frequent code review. Review before commit: All code has at least a cursory peer review before commit. This arises naturally during pairing, but can happen otherwise.

Ensure that VML solutions architects conduct an Architectural Review at the beginning of every project and again just before launch. Long projects may warrant additional review.

Ensure

- Quality is the responsibility of everyone on the team.

Key points

- Before QA, developers verify in approved browsers in the development environment.
- Developers list incomplete or broken items from key documents in QA request.
- Requirements changes appear with screenshots in the test plan.

Demonstrate

Demonstrate your running system often (at least every two weeks). To the final client is better but an internal demo is better than no demo. At the end of the demo, revalidate the plan and budget. PM and Tech Lead are mandatory.

Key Points

- Conduct the demo from a QA or Stage environment (not a development environment)
- Make sure QA has completed a round of testing prior to the demo.
- Do a practice run to see what is working well and what is still in progress.
- Invite AM, PM, & Creative to the practice run so they understand what the client is going to see.
- If you have significant issues, reschedule

Scale

We have more than 50 team leads skilled in this method which means we can quickly divide and conquer to bring timelines in and tackle emergent opportunities as they arise. We've thought a lot about how to create and empower teams so that both experimentation and collective learning grow as your team expands.

Bringing all of this together has enabled DevOps through JDD which has led to dramatic improvements in delivery. At one client, we moved them from 10 releases per year to more than 10 releases per day in production using canary

releases with a dramatic increase in quality: we had a 10 month stretch of production releases with no severity 1 or 2 defects reported (a company best).

Reinvent

Our process is organic and ever-changing; created, modified, and reborn by hundreds of 5-9 person teams over decades delivering software systems.

We talk about this process internally in meetings (and hallways) as well as externally at conferences. In 2019, we had the idea of naming and versioning our process. From that came the idea of open-sourcing the process.

This book is the result. We know this book is on its own journey. We hope to continue to modify it after publication. We hope that you use and modify it as well.

Dedication

To the individuals on our teams defining and redefining Journey-Driven Development with each delivery and retrospective. Though many have contributed topics to this book through numerous commits and pull-requests, even more have contributed ideas.

Errata

We declared a 1.0 a while back (or was it 0.1). There is still so much to do. Adding this page to hold a changelog of sorts for future releases.

Acknowledgements

David Mitchell began this project and is the principal author and editor. Many people contributed content to this book. A comprehensive list can be derived from the git history of our project, but what follows are some specific authors and sections (apologies for any omissions):

1. Chandler, Preston Journey-Driven Agility
2. Heryer, John: Journey-Driven API Design
3. Martin, Jennifer: Maturing the Practice
4. Sykes, Jeff: DevOps

The Pro Git book inspired us to use Pandoc for this.

Discovered after we started that Daniele Dellafore wrote about something similar she called User-Journey Driven Development in August of 2015.

Thanks to Billie Thompson (a.k.a., PurpleBooth) for providing the starter contributing.md template, which we took from Good-CONTRIBUTING.md-template.md.

Process

On some level, this whole book is about process. We added a specific section on process because we believe that methodology concerns itself first with people, process, and tools. Through revisions, we extracted architecture from the tools list and techniques as a sort of smaller scale process. I'm struggling here to avoid calling a technique a micro-process, but it seems clear to us that what remains here after removing techniques is a sort of macro-process, or process in the large.

Journey-Driven Agility

We practice agile software development every day. Fundamentally, we need an approach to generating feedback rapidly and iteratively, repeating these essential steps:

1. Pick a direction to go.
2. Take a step in that direction.
3. Evaluate progress and impact.

How might we embellish this high-level process using a journey-driven approach? Consider an ask from a client or customer. A common origin story is a request from a client to achieve a classic growth curve:




Figure 1: A curved line with one arrowhead showing accelerating growth from right to left, bending upward towards a destination.

Our experience tells us that to get there we need to iterate. This idea of iteration is an ancient idea of infinite rebirth is often represented as an ouroboros, or snake

eating its own tail:




Figure 2: A single arrowhead line wrapped around a circle, pointing back to itself; an ouroboros or snake eating itself.

Segmenting this iteration into delivery and feedback, we've modeled our 3 point list. The lower line is the next step and the feedback line is the evaluation. The checkpoints are the release and the evaluation. Pick a direction, take a step in the direction (delivery), then evaluate progress and impact (feedback) before the next action:

Engineering

It can be this simple. You can layer in a variety of detailed methods. Consider a daily loop using Scrum. The left checkpoint is the daily scrum. The right checkpoint is the daily release. The solid line is the work from the backlog to the next release. The dotted line is the test feedback from the last release (returning us to the next daily scrum).

On the same Scrum team, each of these daily iterations would form a larger loop for the sprint. The left check point is the retrospective from the prior sprint. The right checkpoint is the demonstration at the end of the sprint. The dotted line is the feedback from the release and the work into the next retrospective.

On most projects there are many teams working on different parts of the project. Architectural boundaries make it possible to decompose large engineering projects into many smaller and more manageable teams. Other patterns emerge at scale with specialized teams for automation and developer experience (DevOps).




Figure 3: One circle made from two arrowhead arcs. The bottom solid and the top dotted. At the end of each line is a smaller circle representing a destination or checkpoint.

Design

Outside of engineering there is often a different kind of iterative loop happening on the product design teams. Feedback from launches can lead to a need to reconsider the direction of the product itself. This is often a different team entirely following a different process:

Think of the left circle as the product design process and the right circle as the engineering process. Since our book is concerned primarily with software development, we won't go into a lot of detail about this process.

Initiation

In our work with clients, we actually see an even earlier process that initiates this system of flywheels:

We call this diagram the sideways snowman. The circles increase in size as the initiation team tends to be small, the design team tends to be larger, and the engineering team tends to be the biggest (though none of these are to scale).

At this point, some labels may help:

Note that most projects will iterate differently but the simplest path might be:

1. Originate
2. Restate




Figure 4: Two circles, each made from two arrowhead arcs. Each circle has two checkpoints (though the middle checkpoint is shared).




Figure 5: A sideways snowman made of 3 circles, each made from two arrowhead arcs. Each circle has two checkpoints and two of the checkpoints are shared creating four total checkpoints.

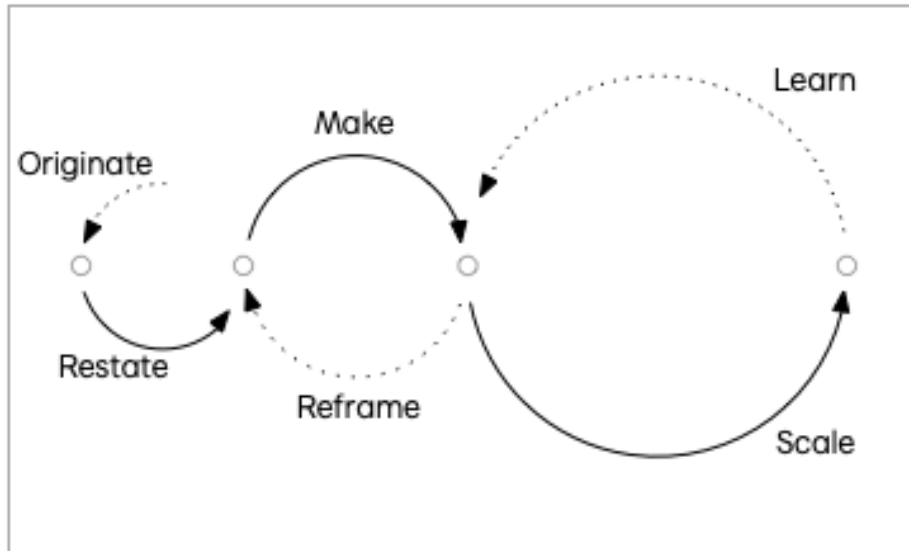


Figure 6: Labeled version of the prior sideways snowman with the six arcs labeled Originate, Restate, Make, Reframe, Scale, and Learn.

3. Make
4. Scale

You might then loop on Learn and Scale for a while before returning to Reframe. The Reframe/Make loop usually runs independent of the Scale/Learn loop. It is even possible that the outcome of Reframing a problem will so reset expectations that you return to Originate though that is infrequent enough that we don't show that as connected.

Adding in the checkpoints, we have this complete diagram:

DevOps

Our DevOps journey at VML emerged as development and operations teams searched for better ways to support the rapid pace of work in an agency. The scripts and manual processes were neither sustainable or scalable enough to deliver as development efforts accelerated. That effort started with destroying the silos that prevented teams from working together. Our collaboration produced automation that enabled teams to build, unit test and deploy code rapidly with greater confidence. As that automation matured, incremental improvements added additional testing from our Quality teams to implement stricter quality gates into application deployment pipelines.

Because DevOps is more about removing silos, we treat automation, scripts and pipelines as tools, not the product of DevOps. We are committed to measuring,

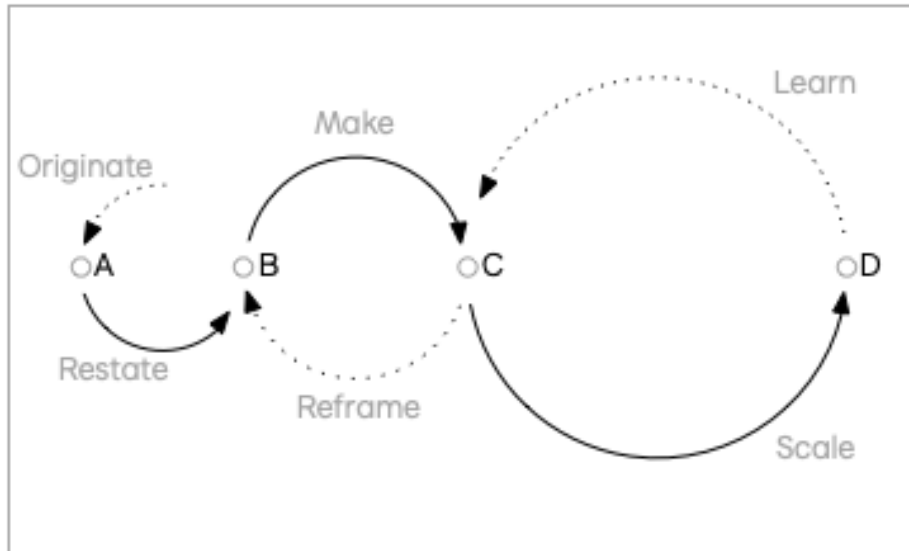


Figure 7: Labeled version of the four checkpoints along with a faded version of the prior labels for the arcs Originate, Restate, Make, Reframe, Scale, and Learn.

refining and improving those tools to improve our application delivery process. We constantly ask the questions: “how might we improve this process?” “What did we learn from this experiment?” “Why did this failure occur?” Asking these questions is the key to learning and continually improving.

When focusing on customer journeys, firms realize that agility and experimentation are key to unlocking value and testing insights. In 2001, the Agile Manifesto articulated a new vision of software development aimed at unlocking value faster. This value creation focused squarely on removing the barriers between business and development in order to deliver software that meets customer needs. Traditional application delivery and deployment techniques seek to improve system stability by slowing the rate of change. By controlling, documenting, and testing changes, organizations attempted to mitigate risk. The competing aims of Agile and preserving system stability lead many to disillusionment with Agile.

The promise of Agile development practices is hampered or lost when following traditional application deployment and delivery techniques. Unlocking the promise of Agile and enabling the kind of experimentation necessary to learn what customers want requires increasing the rate of change not slowing changes down.

DevOps seeks to tear down the barriers between those responsible for developing and operating software. As barriers fall, operational processes must evolve in order to deliver speed and stability simultaneously. Efforts that do not require

both speed and stability as equal goods are sure to deliver neither.

Accounts of DevOps are often told as stories. Perhaps the most famous DevOps story to date is *The Phoenix Project* by Gene Kim.¹ This fictional account tells the tale of Bill Palmer who is elevated lead software delivery for Parts Unlimited. Palmer is coached by a sage advisor who teaches him the Three Ways and helps him understand how lean manufacturing applies to the delivery of IT Systems. These techniques ultimately allow him to help deliver software that allows the business to deliver software to meet customer needs.

The close history of DevOps and Agile is popularly told as Patrick DeBois was frustrated by walls of separation between application development and operations teams.² In 2008, DeBois met with Andrew Schafer to discuss the topic of “Agile Infrastructure.” One of the first public expressions of what “Agile Infrastructure” or “Agile Systems Administration” might look for was presented in 2009 at the O’Reilly Velocity Conference. John Allspaw and Paul Hammond described the process and tools they used to achieve more than ten deploys per day.³

Architecture

API Forward

API Forward is a pattern for evaluating a technical architecture. Why forward:

- forward from behind the CMS
- forward in front of systems of record
- forward in time (like API First) before reference implementation

Forward From Behind the CMS

Content Management Systems (CMS) dominate the technical thinking of marketing teams. We often get feedback when we show diagrams that don’t have the CMS or digital asset management platform (DAM) at the center, coordinating activity.

This makes sense for sites that are purely informational. As sites begin to offer transactions, several patterns emerge:

1. Serverless. Use other people’s services (and servers). Your clients already use other services; you can too. Prefer API Forward for custom services.

¹Gene Kim, Keven Behr, and George Spafford, *The Phoenix Project: A Novel about IT, DevOps, and Helping Your Business Win* (Portland: IT Revolution, 2018).

²Steve Mezak, “The Origins of DevOps: What’s in a Name?”, accessed March 8, 2019, <https://devops.com/the-origins-of-devops-whats-in-a-name/>

³John Allspaw and Paul Hammond, “10+ Deploys Per Day: Dev and Ops Cooperation at Flickr”, accessed March 8, 2019, <https://www.youtube.com/watch?v=LdOe18KhtT4>

2. CMS-hosted services. Host your custom services alongside your CMS' existing "headless" content services. Convenient in the near term. Prefer API Forward as scale or complexity increase.
3. API Forward or something similar with a different name. You may derive API Forward from first principles. Related patterns like Service Facade, Facade and Front Controller share some structure and consequences. Consider API Forward as a framework to evaluate your solution's fit for modern marketing.

In short, consider API Forward as you move from content-only websites. Specifically as you add commerce, other transactions, or even as you add things like chatbots as conversational state and transactional state are analogous.

Forward of Systems of Record

We've also seen several anti-patterns emerge:

1. CMS at the center. This is an unintentional mediator that comes from starting with a content platform and building everything around it. Unless you are a traditional publisher, the content platform should not be the center of your architectural diagram. Use API Forward to create an effective mediator connecting you to your customer's needs.
2. Stranger Danger. This happens when you don't have any mediator and can arise from simply having line of business systems provide UI directly to your customers. Sometimes referred to as violations of the what Design Patterns referred to as the Law of Demeter or more simply, Don't Talk to Strangers. In an architectural diagram, you see this as skipping cake layers in the diagram. Use API Forward to create more reusable architectural components that aren't tied to either back-end platforms or front-end constraints.

Forward in Time

When developing an API, several patterns emerge:

1. Implementation First (a.k.a Reference Implementation). The implementation is the specification. API Forward recommends coding against a specification, which sounds more like...
2. API First (a.k.a. Contract First, Design By Contract). The implementation must meet the specification. API Forward leverages the technical approach API First in that we sketch a front-end UI design to translate the end-user's needs to an API endpoint.
3. API Forward or something similar. Methods like Responsibility-driven design (RDD) and Domain-Driven design (in particular the idea of Ubiquitous Language), may yield an API similar to one aligned with API Forward. Use API Forward as an iterative approach aligned with modern marketing.

When: Adding transactions. Also consider in conversational commerce, as conversations are transactional as well.

Uses Customer Language

API Forward is journey-driven. That is, it uses the the language of the customer, not the language of enterprise components; you likely have existing API endpoints that may not be written in customer language. We do not expect to rewrite them. Instead, API Forward tends to lead to a layer of orchestration aggregating and sequencing lower-level APIs. API Forward lets you leverage existing and future enterprise investment.

API Forward and Microservices

Microservices.io provides a good baseline for understanding what we mean by a microservice:

1. Highly maintainable and testable
2. Loosely coupled
3. Independently deployable
4. Organized around business capabilities
5. Owned by a small team

Most modern discussions of microservices include a few more expectations, such as:

1. Microservices execute and participate in a service mesh
2. There is a consistent and predictable contract across services, at a macro-level
3. Are easily composable with other services
4. Fine grained, distributed traceability
5. Graceful fault tolerance and resiliency

Unpacking the service mesh a bit more, we find these additional expectations:

1. Dynamic registration/discovery
2. Independent and precise auto-scaling
3. Health check telemetry
4. Polyglot support

Content Rocket

For brands, personalization at scale is the marketing holy grail. But when people try to create these experiences, it is easy to get stuck in silos: attempting to deliver a personalized experience using disconnected technology. Personalization can become more about internal turn wars than anticipating customer needs.

Content Rocket is an end-to-end framework that puts the customer at the heart of our efforts and illustrates how we deliver personalized experiences at scale.

Data

Everything starts with a customer, which is to say, everything starts with customer data. Start by gathering qualitative and quantitative insights using both first and third party data to better understand our customers, their needs, and their journey.

Strategy

In a content process named Content Rocket, you might be surprised that we also consider experience strategy in concert with content strategy to uncover the best use of content in context of the journey. Experience can be a fundamental unlock for putting content into this context.

1. Experience Strategy
2. Content Strategy

Experience Strategy

1. Personas
2. Jobs Theory
3. KPIs
4. Value-centered design
5. Think-Make

Content Strategy

1. Editorial
2. Product
3. Campaign
4. Offer
5. Utility

Asset Creation

1. Audience story
2. Core assets
3. Copywriting
4. Proofing
5. Editing
6. Translation
7. Legal review
8. Tagging

Activation

Primary tasks:

1. Identify

2. Assemble
3. Serve

Platforms:

1. Digital Asset Management
2. Content Management
3. Customer Relationship Management
4. Customer Data Platform / Data Management Platforms
5. Dynamic Creative Optimization
6. Artificial Intelligence / Machine Learning

Journey Audience Map

Evolution

Reuse

Refine

Reevaluate

Techniques

Journey-Driven API Design

Journey-Driven API Design is a contract-first approach to API creation. Contract-first means the implementation must meet a specification. The result is customer value alignment with technology through a well documented and designed API that services multi-channel client interactions.

Journey-Driven API design consists of three primary phases of analysis:

- Domain
- System
- Persistence

The activities in each phase are ordered and should be worked continuously and iteratively throughout a project. Each step should have a design artifact for documentation, typically a technical diagram. Document significant architectural decisions as an Architectural decision Record.

Document the decisions

A technical decision may be “architecturally significant” due to its potential for long-lasting impact. Context is lost as projects evolve, and it is vital to capture the architectural decisions made. Preserve essential aspects through the design process with Architecture Decision Records.

Architecture Decision Record Capture architecture decisions with an Architecture Decision Record. The format consists of a minimal set of information needed to understand: factors that influence a decision specific items decided anticipated outcomes of the decision

Resource Analysis

Resource Analysis provides a technical architect the inputs required to bridge technology and customer value. These activities focus on understanding the problem domain and requirements to guide API design.

1. Gather Functional Requirements Understand the problem domain by working with business stakeholders and product owners to learn how customers will interact with the system. The Functional Requirements Document (FRD), User Stories, Statements of Work, and even UI/UX inputs can help inform the API design. Common scenarios included customers interacting with a web or mobile application supported by APIs.

2. Gather Non Functional Requirements Non-functional requirements (NFR) are the constraints on operational behavior. Information such as response time requirements, performance budgets, accessibility requirements, legal and regulatory compliance constraints are inputs into a technical system design. These constraining attributes need to be realized early and designed into the architectural foundation of the system. Work with business and product owners to define non-functional requirements.

3. Create a Domain Model (Conceptual ER Model) A domain model is a living document, updated as requirements evolve. It establishes a ubiquitous language with domain experts and identifies fundamental entity types, behaviors, and relationships. Start a domain model by extracting key nouns and verbs from user stories or requirements. The domain model is a prerequisite to creating the API Contract.

4. Design an API Contract Create the API contract first, so the implementation meets the specification. The OpenAPI spec is an Interface Description Language (IDL) to design and document RESTful APIs. The definition of an API with an IDL provides clear direction and acceptance criteria to API developers and consumers by way of a well-defined contract. The document should be tracked in source control with the code and updated as the API evolves.

System Analysis

System Analysis is valuable to technical and non-technical audiences. Analysis with diagrams captures the current and future state of systems, including their relationships at several abstraction levels. Start with the highest level that delivers value by identifying who is using which system. Zoom into individual

technical systems to visualize the deployable containers that represent application or data. Finally, capture an understanding of how application functionality is related and encapsulated down to the source code representation.

Describing how the API will fit into a new or existing system can be difficult. The C4 Model diagramming framework helps technical and non-technical audiences navigate and communicate complex software systems. The following steps describe the parts of defining and documenting a technical system using C4.

1. Create a System Context Diagram A system context diagram provides a starting point, showing how the software system in scope fits into the world around it. This diagram describes the type of software system, who uses it, and how it fits into an existing environment.

2. Create a Container Diagram A container diagram zooms into the software system in scope, showing the high-level technical building blocks (containers) and how they interact. This diagram shows the high-level technology decisions of the architecture and how it distributes responsibility. Additionally, it illustrates the relationships that containers have to each other.

3. Create a Component Diagram A component diagram represents the pieces into an individual container. It demonstrates the high-level structure of an application by describing the responsibilities of components and how they interact.

4. Create a Code Diagram UML Class diagrams and package diagrams describe how to implement a component as code. There is value in proactively creating UML Class and Package diagrams to communicate important software design concepts and patterns.

5. Document Sequences and Interactions Capture interactions between components. Sequence diagrams visualize the message flow as a communication tool to describe complex interactions over time. Documenting advanced interaction paths clarifies processes to non-technical audiences and informs acceptance criteria used by development teams.

Persistence Analysis

If the API needs persistence, create a logical and physical data model. Use the domain model or conceptual model from the Resource Analysis phase as an input.

1. Create Entity Relationship Diagram (Logical Model) Expand upon the domain model and create an (ERD) as a logical representation of the data entity properties and relationships.

2. Create Database Schemas (Physical Model) Build on the ERD created and database schema. If an RDBMS is needed, a schema is required. Design a schema for systems using non-relational data storage.

3. Create a caching strategy Data access patterns and available technologies should drive the definition of a caching strategy. The performance-focused NFRs should be an input to the caching strategy. Identify available cache solutions, such as Akamai and Redis, to address data at different architecture levels. Describe a caching strategy by detailing use-case scenarios around data access, caching, and eviction policies.

Maturing the Practice

Maturing a practice is a battle against decay and entropy. You require an infusion of energy and purpose; maturation never happens by default. High performing development teams hold a commitment to maturation because the quality of the software depends on it.

Evolving a team starts with a decision from somewhere in the hierarchy. Identifying a path for affecting those changes can be as challenging setting the direction.

Approach the process with a combination of directive and suggestive measures to allow the practice to fulfill software quality expectations with fully invested developers who embrace team values and demonstrate ownership of results.

Directive vs Suggestive Approaches

The *directive* approach hinges on the authority to enforce a mandate and is effective when the team has little option beyond compliance. For example:

- Edict 1: PR's must have two reviewers before being merged
- Enforcement: Configure github project to require 2 approvals before merging
- Result: All merge PR's have received 2 approvals

Mission Accomplished. When the letter of the law is all that is required and when the decision maker has the authority to compel compliance, this is a simple, direct path.

The vulnerability of the directive approach surfaces when the decision maker is looking to change internalized values.

- Edict 2: Developers resolve linting violations before merging a PR
- Enforcement: Configure github project to block merging if there are linting violations beyond a reasonable threshold
- Result: Some developers resolve the violations; other developers raise the violation threshold

Mission Failed. At this point, enforcement becomes a bit more harsh and possibly personal. It risks affecting the cohesiveness of the team. What is needed here is for team members to embrace the value of minimizing linting violations and to own the practice of resolving them immediately. Suggestive tactics are needed to fill the gap.

The *suggestive* approach works like harvesting a crop. There is soil to prepare, as well as seeds to plant. Then, with purposeful cultivation, there will be a crop to harvest: At an abstract level, it looks like:

- prepare the soil -> build trust and establish credibility,
- plant seeds -> sell the value of proposed evolution, and
- cultivate the crop -> capitalize on timely opportunities.

A suggestive approach takes time and patience. Additionally, there is a synergistic benefit of the rich soil of trust and credibility being available for sowing other seeds. In reality, most attempts to mature a practice will benefit from a combination of both approaches.

Example: Increase Unit Testing Adoption

You want to increase adoption of unit testing. First consider evolving your internal practice where you have authority to set a directive and then consider evolving your client's practice where you might not be able to set a directive.

Evolve Our Practice How might we evolve our practice to include more unit testing. Consider three approaches:

1. Strictly directive
2. Strictly suggestive
3. Suggestive with directive

Strictly Directive: Issue an edict: Reject all PR's without unit tests.

- Many developers feel resentment at being forced into a practice they will argue is not valuable
- A few developers of a certain mindset accept the challenge and become good testers
- Others get really good at finding ways around testing
- Still others leave the team

This brute force way produces some testers, but it very likely poisons the environment for future changes and results in practices that evade the mandate without fulfilling it.

Strictly Suggestive: make team members feel safe, convince the team of the value of testing, remove obstacles to testing

- Probably fewer feelings of resentment
- More developers, but not all, become good testers

- Fewer developers, but not none, find ways around testing
- Likely none will leave

Because resistant team members avoid adopting the practice, this approach is not successful.

Suggestive with Directive: use the suggestive approach along with rejecting PR's without unit tests

- Dramatically less resentment
- Dramatically more good unit testers
- Very limited avoidance or attrition

Eventually, we adopt the practice fully.

Evolve Our Client's Practice to Include Unit Testing The twist here is that we lack direct authority to use the directive approach within a client organization. We begin with suggestive strategies to influence developers to embrace unit testing and to influence client leadership to eventually include directive strategies.

Turns out suggestive fluency is a super power. A confluence of micro actions point to the overall objective. Every context requires its own implementation details. Here are some ideas to start with.

1. Prepare the soil
2. Plant the seeds
3. Cultivate: Capitalize on timely opportunities

Prepare the Soil We have 3 goals when preparing the soil:

1. Build trust
2. Establish credibility
3. Avoid harming the environment

Build Trust

Make it abundantly clear that our goal is to make the client's life easier and project better

Demonstrate an urgency to relieve client devs' pressure points

Work on personal connections with all client devs

Establish Credibility

Walk the walk. Include unit tests in all PR's, with exceptionally rare exceptions

Drop everything if someone wants to pair. Drop everything faster if someone wants to pair on testing<

Demonstrate a willingness to do the heavy lifting of establishing a practice of testing

Communicate personal story of learning to unit

Avoid Harming Environment

Do nothing that could possibly communicate an attitude of superiority

Never diminish the objections to the challenges of testing

Readily acknowledge the learning curve for testing is an investment

Avoid sounding righteous about testing. A good technique is to “blame” VML leadership for the requirement to test: “We’re graded on testing our code”

Plant the Seeds We have 3 goals when planting the seeds:

1. Testing is beneficial for developers
2. Developers Own the Code Quality/Success
3. Unit Testing is inevitable

Testing is beneficial for developers

Identify and respond to the primary objections to testing by facilitating group discussions

Developers Own the Code Quality/Success

Make a concerted effort to replace 3rd person pronouns with 1st person pronouns:

- THEY should -> WE should
- The Business wants -> WE want
- QA’s bugs -> OUR bugs

Move from “we are victims of this mess” to “we are mastering this domain”

Unit Testing is inevitable

Regularly make statements that assume the change will be implemented:

“When we have better test coverage, we won’t have so many bug fixes regressing”

“When we have better test coverage, we’ll be able to make these changes easier”

“Your job is going to be so much easier when our coverage improves”

Cultivate: Capitalize on Timely Opportunities We have four opportunities to capitalize on when cultivating:

1. Retros
2. Light Sprints
3. New Hires
4. PR Reviews

Retros

Since retros produce action items, this is a great place to talk about anything related to lack of testing.

Celebrate improvements in testing.

Light Sprints

Identify tech debt that could promote unit testing. Get tickets created, and lobby for that work to be included in sprints.

New Hires

Develop allies by onboarding new client hires and incorporating pairing and testing into their orientation.

PR Reviews

In the comments of PR's without tests, provide code snippets for adding tests or offer to pair for writing tests to cover the work.

Withhold PR approval if there are no tests. Important to read the room here and time this well. Under certain conditions, this could be perceived as aggressive and do more harm than good.

Pattern Languages

People

Journey of the Software Professional

Manager Tools Adoption

Evolution of Teamwork

People Matter

See also:

1. Peopleware: Productive Projects and Teams
2. Multipliers Book
3. Radical Candor: Be a Kickass Boss | The Book

Tools

Future Topics

Agile Release Planning

Answers the question: How do you estimate project costs, client pricing and delivery timelines? This is a favored method for coming up with an agile release plan. We have used this method successfully on a number of projects. We adapted this approach from training by Gregory Smith of the Agile Alliance. Once we have a story backlog, we move through a series of steps to create a release plan that maps out the stories into sprints. The plan addresses how long the project is estimated to take, and how many points per sprint the team can take on. We take the anticipated staffing and development cadence along with the timeline and derive project costs.

STEP 1: ESTABLISH THE TARGET VELOCITY

First we take the stories and lay them all out on a big table. We start with wild guess groupings. No scale or points are disclosed to the team as a basis for evaluating relativity.

The point here is to establish relativity to one another in size groupings, without indicating precision. For example, some may use a scale of sm, med, lg, xl, xxl . But we need to steer clear of qualifying a medium as “exactly twice as big as a small,” or a large is “three times a small,” etc. We seek rough sizing relative to one another. Med is simply more complex than a small, large more complex than med. To emphasize relativity and remove any precision in the scale (and have a little fun), VML favors “fruit sizing” using a planning poker deck consisting of grape, orange, coconut, pineapple, watermelon.

1.1 CALIBRATE STORIES

After assigning relative sizes to stories, the team goes through an exercise in calibration of the sizes. They group by size and look across the group for consistency. Are all grapes really grapes, or are there some oranges hiding among grapes? We perform boundary analysis — look for the biggest small and the smallest medium and evaluate that the stories are in the right size group.

1.2 EVALUATE DISTRIBUTION

Another thing we do is plot the distribution, or the number of stories in each category to evaluate the distribution of stories. Sometimes if the team lacks enough criteria to scale stories or the scale itself is not understood, a trend may arise where “everything’s an orange” which is problematic, and shows the team needs to talk more or the scale may need to be addressed. If everything were XL that would be a similar problem too.

1.3 RUN SCENARIOS OF SIZINGS IN SPRINTS

Next the team goes through an exercise of looking at each category and estimating the number of stories per sprint (# per sprint) they can fit in. Assuming a two-week sprint, the team grabs as many smalls as they think they can fit in a sprint. Cycle through this three or four times for smalls, then move on to mediums and do the same. Repeat for the rest of the sizes and record the findings 18 (scenarios a, b, c, d).

1.4 ESTABLISH TOTAL SCOPE

We pick a scale to assign numerical weight values to the categories. A scale based on the Fibonacci sequence is used most often (grape = 2, strawberry = 3, orange = 5, pineapple = 8, watermelon = 13). The first step toward total scope is to calculate the value per fruit category by multiplying the number of stories by the corresponding numerical weighting. Total the values across the row and this value becomes your total scope.

Code Review

Pull Request

Pairing as Review

Pride and Shame

Like, Wish, Wonder

Turn Watermelons Into Orangeberries

On a recent transformation project, we used Fruit Sizing with Planning Poker decks.

A watermelon is a big story, effectively a marker to remind you to break it into smaller pieces. That sounds like work, slicing the watermelon into bit sized chunks. But a technology optimist sees a chance for transformation.

As is often the case, I draw my inspiration from Kent Beck, his eXtreme Programming process, specifically his admonition to DoTheSimplestThingThatCouldPossiblyWork. Question your assumptions.

Back to the work...

For the purposes of Agile Release Planning, our teams estimate their cumulative scope distance across 6 iterations is around 600 points (+/- 30%). That's just a sizing estimate but it helps product owners set priority. In this case, the team is looking at a 1200 point wish list. Time to make some cuts.

Fortunately, they had 10 watermelons (around 500 points).

While hanging out in one of the API rooms when a product owner came looking to talk with a pod lead about a watermelon. One goal is to reduce calls to the customer care centers. A big cause of calls is bill shock. Any change in the bill

can trigger a call. Also, it turns out some people are more likely than others to call.

At first blush, sounds like a watermelon. But that depends on how you slice it:

1. Amount of change rule. Easy. Compare bill A to bill B and see if it exceeds a tolerance. If it isn't big, don't send.
2. Likely to call rule. Hard. Doing this well requires that we create a predictive model of customer behavior. Classic big data problem.
3. Do not contact rule. Easy. Company had this exposed as an API. Reuse.
4. Actually sending a notification. Easy. Company already had an API for contact. Reuse.

So, the size is all in that second slice. We spent some time imagining lots of ways how you could determine who would call. Undefined things sound hard. The product owner proposed something simple, "What if I can make it easy for the initial Minimum Viable Product (MVP) release? What if it was just a text file?"

The pod lead understood where these things come back to bite you, "How often would it need to change? How will we manage the changes to the file?"

"Assume it doesn't change for the MVP scope."

"Is there any PII data in the list?"

"We could use the this internal account number, that's not PII."

"Well, then it's easy."

"How easy. Grape, strawberry, or orange?"

We settled on orangeberry.

Back Matter

Bibliography